

Recommendations with IBM

In this notebook, you will be putting your recommendation skills to use on real data from the IBM Watson Studio platform.

You may either submit your notebook through the workspace here, or you may work from your local machine and submit through the next page. Either way assure that your code passes the project [RUBRIC](#). **Please save regularly.**

By following the table of contents, you will build out a number of different methods for making recommendations that can be used for different situations.

Table of Contents

I. Exploratory Data Analysis

II. Rank Based Recommendations

III. User-User Based Collaborative Filtering

IV. Content Based Recommendations (EXTRA - NOT REQUIRED) V. Matrix Factorization

VI. Extras & Concluding

At the end of the notebook, you will find directions for how to submit your work. Let's get started by importing the necessary libraries and reading in the data.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import project_tests as t
import pickle
```

```
%matplotlib inline
```

```
df = pd.read_csv('data/user-item-interactions.csv')
df_content = pd.read_csv('data/articles_community.csv')
del df['Unnamed: 0']
del df_content['Unnamed: 0']
```

```
# Show df to get an idea of the data
df.sort_values(by="article_id").head(5)
```

```
      article_id
title \
1935      0.0  detect malfunctioning iot sensors with streami...
40755      0.0  detect malfunctioning iot sensors with streami...
```

```

42948      0.0  detect malfunctioning iot sensors with streami...
2328      0.0  detect malfunctioning iot sensors with streami...
19242      0.0  detect malfunctioning iot sensors with streami...

```

```

                                email
1935  384255292a8223e84f05cale1deaa450c993e148
40755  8bd0afc488016810c287ac4ec844895d570b0af4
42948  451a9a4a4cb1cc4e5f38d04e8859cc3fb275cc66
2328   ca7d48adf2c7394ed5a8776de959fa8047e43d4b
19242  a60b7e945a8f2114d5dfbdd53182ad1d526534e2

```

```

# Show df_content to get an idea of the data
df_content.head()

```

```

                                doc_body \
0  Skip navigation Sign in SearchLoading...\r\n\r...
1  No Free Hunch Navigation * kaggle.com\r\n\r\n ...
2  ≡ * Login\r\n * Sign Up\r\n\r\n * Learning Pat...
3  DATALAYER: HIGH THROUGHPUT, LOW LATENCY AT SCA...
4  Skip navigation Sign in SearchLoading...\r\n\r...

```

```

                                doc_description \
0  Detect bad readings in real time using Python ...
1  See the forest, see the trees. Here lies the c...
2  Here's this week's news in Data Science and Bi...
3  Learn how distributed DBs solve the problem of...
4  This video demonstrates the power of IBM DataS...

```

```

                                doc_full_name doc_status
article_id
0  Detect Malfunctioning IoT Sensors with Streami...      Live
0
1  Communicating data science: A guide to present...      Live
1
2          This Week in Data Science (April 18, 2017)      Live
2
3  DataLayer Conference: Boost the performance of...      Live
3
4          Analyze NY Restaurant data using Spark in DSX      Live
4

```

Part I : Exploratory Data Analysis

Use the dictionary and cells below to provide some insight into the descriptive statistics of the data.

1. What is the distribution of how many articles a user interacts with in the dataset? Provide a visual and descriptive statistics to assist with giving a look at the number of times each user interacts with an article.

```
# To find the number of articles an user has interacted with, lets first group the "user-item-interactions" dataset by the "email" column
```

```
groupby_email = df.groupby(by="email")
```

```
# check to see if groupby_email DataframeGroupBy object is instantiated
```

```
groupby_email
```

```
<pandas.core.groupby.generic.DataFrameGroupBy object at 0x000002C048DF7640>
```

```
# Lets now see the number of articles each user has interacted with  
user_interaction_counts=groupby_email["article_id"].count()  
user_interaction_counts
```

```
email  
0000b6387a0366322d7fbfc6434af145adf7fed1    13  
001055fc0bb67f71e8fa17002342b256a30254cd     4  
00148e4911c7e04eef8def7bbbdaf1c59c2c621     3  
001a852ecbd6cc12ab77a785efa137b2646505fe     6  
001fc95b90da5c3cb12c501d201a915e4f093290     2  
..  
ffc6cfa435937ca0df967b44e9178439d04e3537     2  
ffc96f8fbb35aac4cb0029332b0fc78e7766bb5d     4  
ffe3d0543c9046d35c2ee3724ea9d774dff98a32    32  
fff9fc3ec67bd18ed57a34ed1e67410942c4cd81    10  
fffb93a166547448a0ff0232558118d59395fecb    13  
Name: article_id, Length: 5148, dtype: int64
```

```
# Fill in the median and maximum number of user_article interactions below
```

```
median_val = user_interaction_counts.median() #50% of individuals interact with 3 number of articles or fewer.  
max_views_by_user = user_interaction_counts.max() # The maximum number of user-article interactions by any 1 user is 364.  
median_val,max_views_by_user
```

```
(3.0, 364)
```

```
# lets now look at some stats around the counts of articles interaction by the users using the describe() method  
user_interaction_counts.describe()
```

```
count    5148.000000  
mean       8.930847  
std       16.802267
```

```
min          1.000000
25%          1.000000
50%          3.000000
75%          9.000000
max         364.000000
Name: article_id, dtype: float64
```

looks like the the max number of articles any 1 user has interacted with is 364
and the median is 3, and on average an user has interacted with around 9 articles

```
# lets create a function to plot a simple visual that shows the  
frequency distribution of user interactions with the articles
```

```
from matplotlib.pyplot import title
```


```
def plot_user_interactions_counts(df,plot_type):
    """
    A function to plot a visual around the user interaction counts  
with the articles

    Paramters
    df: name of the df that is grouped by email and has counts of  
articles per email
    plot_type: kind of chart in a string format

    Returns
    Plot with frequency of user interaction with the articles
    """

    return df.plot(kind=plot_type, title="user-article interactions")

# Calling the plot_user_interactions_counts function to plot the user  
interaction counts using a histogram
plot_user_interactions_counts(user_interaction_counts,"hist");
```



the above histogram shows the distribution of interactions

2. Explore and remove duplicate articles from the **df_content** dataframe.

Find and explore duplicate articles

lets first look at a sample of the dataframe

`df_content.sample(5)`

	doc_body \	doc_description \	doc_full_name	doc_status
805	CONNECTING TO COMPOSE FOR MYSQL\r\nShare on Tw...	Here, we'll take you from provisioning your ow...	Connecting to Compose for MySQL	Live
49	Compose The Compose logo Articles Sign in Free...	We'll also look at using PostGIS to filter our...		
502	G. Adam Cox Blocked Unblock Follow Following J...	The SETI@IBMCloud project and the SETI Institu...	GeoFile: Using OpenStreetMap Data in Compose P...	Live
3	DATALAYER: HIGH THROUGHPUT, LOW LATENCY AT SCA...	Learn how distributed DBs solve the problem of...		
386	Skip to main content IBM developerWorks / Deve...	See how to load new geospatial data into your ...	Defensive IBM Object Storage Containers – IBM ...	Live

```

501
3    DataLayer Conference: Boost the performance of...      Live
3
386  Load geospatial data into dashDB to analyze in...      Live
385

```

```

# lets first check for any duplicate rows as whole
df_content.duplicated().sum()

```

```
0
```

there are no duplicate rows as a whole (repeated rows)

```

# lets check for duplicate article ids
df_content[df_content.duplicated(subset=["article_id"])]

```

```

doc_body \
365 Follow Sign in / Sign up Home About Insight Da...
692 Homepage Follow Sign in / Sign up Homepage * H...
761 Homepage Follow Sign in Get started Homepage *...
970 This video shows you how to construct queries ...
971 Homepage Follow Sign in Get started * Home\r\n...

```

```

doc_description \
365 During the seven-week Insight Data Engineering...
692 One of the earliest documented catalogs was co...
761 Today's world of data science leverages data f...
970 This video shows you how to construct queries ...
971 If you are like most data scientists, you are ...

```

```

doc_full_name doc_status
article_id
365          Graph-based machine learning      Live
50
692 How smart catalogs can turn the big data flood...      Live
221
761 Using Apache Spark as a parallel processing fr...      Live
398
970          Use the Primary Index      Live
577
971 Self-service data preparation with IBM Data Re...      Live
232

```

```

# Lets take an article Id from the above results and have a look at
the duplicate rows with that article Id
df_content[df_content.loc[:, "article_id"] == 221]

```

```

doc_body \
221 * United States\r\n\r\nIBM® * Site map\r\n\r\n\r\n...
692 Homepage Follow Sign in / Sign up Homepage * H...

```

```

doc_description \

```

```
221 When used to make sense of huge amounts of con...
692 One of the earliest documented catalogs was co...
```

```

                                doc_full_name doc_status
article_id
221 How smart catalogs can turn the big data flood...      Live
221
692 How smart catalogs can turn the big data flood...      Live
221
```

```
# Remove any rows that have the same article_id - only keep the first
df_content = df_content.drop_duplicates(subset=["article_id"],
keep='first')
```

```
# Test if the drop duplicate method worked by pulling an article id
from one of the previous results (pre-drop)
df_content[df_content.loc[:, "article_id"] == 221]
```

```

                                doc_body \
221 * United States\r\n\r\nIBM® * Site map\r\n\r\n...
```

```

                                doc_description \
221 When used to make sense of huge amounts of con...
```

```

                                doc_full_name doc_status
article_id
221 How smart catalogs can turn the big data flood...      Live
221
```

Looks like the drop_duplicate method worked fine

3. Use the cells below to find:

- a.** The number of unique articles that have an interaction with a user.
- b.** The number of unique articles in the dataset (whether they have any interactions or not).
- c.** The number of unique users in the dataset. (excluding null values)
- d.** The number of user-article interactions in the dataset.

```
# to look at the number of unique articles that have an interaction
with a user , lets first look at articles in the df
df["article_id"]
```

```

0      1430.0
1      1314.0
2      1429.0
3      1338.0
4      1276.0
...
45988  1324.0
45989   142.0
45990   233.0
```

```
45991    1160.0
45992     16.0
Name: article_id, Length: 45993, dtype: float64
```

```
# lets do a unique value count using the nunique() method on the above series to find out the number of unique articles that have an interaction with a user
df["article_id"].nunique()
```

```
714
```

```
# to find the number of unique articles in the dataset (whether they have any interactions or not). Lets look at the articles in df_content and use the nunique() method
df_content["article_id"].nunique()
```

```
1051
```

```
# to find number of unique users in the dataset. (excluding null values), lets just call the nunique() method on the email column of the df dataframe
df["email"].nunique()
```

```
5148
```

```
# to find the number of user-article interactions in the dataset, lets just call the shape attribute on the df dataframe and use slice notation to find the number of rows
df.shape[0]
```

```
45993
```

```
unique_articles = df["article_id"].nunique() # The number of unique articles that have at least one interaction
total_articles = df_content["article_id"].nunique() # The number of unique articles on the IBM platform
unique_users = df["email"].nunique() # The number of unique users
user_article_interactions = df.shape[0] # The number of user-article interactions
```

4. Use the cells below to find the most viewed **article_id**, as well as how often it was viewed. After talking to the company leaders, the `email_mapper` function was deemed a reasonable way to map users to ids. There were a small number of null values, and it was found that all of these null values likely belonged to a single user (which is how they are stored using the function below).

```
# to find the most viewed article_id, as well as how often it was viewed.
# lets call the value_counts() method on the article_id column of the df dataframe
# and also use the head(1) method to display the top most article_id
df["article_id"].value_counts(ascending=False).head(1)
```



```
1429.0    937
Name: article_id, dtype: int64
```

```
most_viewed_article_id = '1429.0' # The most viewed article in the
dataset as a string with one value following the decimal
max_views = 937 # The most viewed article in the dataset was viewed
how many times?
```

```
## No need to change the code here - this will be helpful for later
parts of the notebook
# Run this cell to map the user email to a user_id column and remove
the email column
```

```
def email_mapper():
    coded_dict = dict()
    cter = 1
    email_encoded = []

    for val in df['email']:
        if val not in coded_dict:
            coded_dict[val] = cter
            cter+=1

    email_encoded.append(coded_dict[val])
    return email_encoded
```

```
email_encoded = email_mapper()
del df['email']
df['user_id'] = email_encoded
```

```
# show header
df.head()
```

	article_id	title
0	1430.0	using pixiedust for fast, flexible, and easier...
1	1314.0	healthcare python streaming application demo
2	1429.0	use deep learning for image classification
3	1338.0	ml optimization using cognitive assistant
4	1276.0	deploy your python model as a restful api
5		

```
## If you stored all your results in the variable names above,
## you shouldn't need to change anything in this cell
```

```
sol_1_dict = {
    '50% of individuals have ____ or fewer interactions.':
```

```

median_val,
    ``The total number of user-article interactions in the dataset is
_____.``: user_article_interactions,
    ``The maximum number of user-article interactions by any 1 user is
_____.``: max_views_by_user,
    ``The most viewed article in the dataset was viewed _____
times.``: max_views,
    ``The article_id of the most viewed article is _____.``:
most_viewed_article_id,
    ``The number of unique articles that have at least 1 rating
_____.``: unique_articles,
    ``The number of unique users in the dataset is _____``:
unique_users,
    ``The number of unique articles on the IBM platform``:
total_articles
}

```

Test your dictionary against the solution

```
t.sol_1_test(sol_1_dict)
```

It looks like you have everything right here! Nice job!

Part II: Rank-Based Recommendations

Unlike in the earlier lessons, we don't actually have ratings for whether a user liked an article or not. We only know that a user has interacted with an article. In these cases, the popularity of an article can really only be based on how often an article was interacted with.

1. Fill in the function below to return the **n** top articles ordered with most interactions as the top. Test your function using the tests below.

```

def get_top_articles(n, df=df):
    """
    INPUT:
    n - (int) the number of top articles to return
    df - (pandas dataframe) df as defined at the top of the notebook

    OUTPUT:
    top_articles - (list) A list of the top 'n' article titles

    """
    # Your code here
    top_articles=list(df.title.value_counts().head(n).index) #
    value_counts() method to create series of counts of titles and head()
    function to retrieve the top n records
    return top_articles # Return the top article titles from df (not
    df_content)

def get_top_article_ids(n, df=df):
    """

```

INPUT:

n - (int) the number of top articles to return

df - (pandas dataframe) df as defined at the top of the notebook

OUTPUT:

top_articles - (list) A list of the top 'n' article titles

'''

Your code here

*top_articles=list(df.article_id.value_counts().head(n).index) #
value_counts() method to create series of counts of ids and head()
function to retrieve the top n records*

top_articles = [str(x) for x in top_articles]

return top_articles # Return the top article ids

`print(get_top_articles(10))`

`print(get_top_article_ids(10))`

`['use deep learning for image classification', 'insights from new york
car accident reports', 'visualize car data with brunel', 'use xgboost,
scikit-learn & ibm watson machine learning apis', 'predicting churn
with the spss random tree algorithm', 'healthcare python streaming
application demo', 'finding optimal locations of new store using
decision optimization', 'apache spark lab, part 1: basic concepts',
'analyze energy consumption in buildings', 'gosales transactions for
logistic regression model']`

`['1429.0', '1330.0', '1431.0', '1427.0', '1364.0', '1314.0', '1293.0',
'1170.0', '1162.0', '1304.0']`

Test your function by returning the top 5, 10, and 20 articles

`top_5 = get_top_articles(5)`

`top_10 = get_top_articles(10)`

`top_20 = get_top_articles(20)`

Test each of your three lists from above

`t.sol_2_test(get_top_articles)`

Your top_5 looks like the solution list! Nice job.

Your top_10 looks like the solution list! Nice job.

Your top_20 looks like the solution list! Nice job.

Part III: User-User Based Collaborative Filtering

1. Use the function below to reformat the **df** dataframe to be shaped with users as the rows and articles as the columns.

- Each **user** should only appear in each **row** once.
- Each **article** should only show up in one **column**.

- **If a user has interacted with an article, then place a 1 where the user-row meets for that article-column.** It does not matter how many times a user has interacted with the article, all entries where a user has interacted with an article should be a 1.
- **If a user has not interacted with an item, then place a zero where the user-row meets for that article-column.**

Use the tests to make sure the basic structure of your matrix matches what is expected by the solution.

create the user-article matrix with 1's and 0's

```
def create_user_item_matrix(df):
    """
    INPUT:
    df - pandas dataframe with article_id, title, user_id columns

    OUTPUT:
    user_item - user item matrix

    Description:
    Return a matrix with user ids as rows and article ids on the
    columns with 1 values where a user interacted with
    an article and a 0 otherwise
    """

    unstacked=df.groupby(['user_id',
    'article_id']).count().unstack().fillna(0) # group by user id and
    article ids and perform count agg and unstack rows to columns while
    filling nulls with 0s
    unstacked.columns= unstacked.columns.droplevel() # drop the extra
    column level (level not needed for our matrix)
    user_item=unstacked.applymap(lambda x: 1 if x > 1 else x ) #
    making sure all the values inside the matrix is either 0 or 1

    return user_item # return the user_item matrix

user_item = create_user_item_matrix(df)

## Tests: You should just need to run this cell. Don't change the
code.
assert user_item.shape[0] == 5149, "Oops! The number of users in the
user-article matrix doesn't look right."
assert user_item.shape[1] == 714, "Oops! The number of articles in
the user-article matrix doesn't look right."
assert user_item.sum(axis=1)[1] == 36, "Oops! The number of articles
seen by user 1 doesn't look right."
print("You have passed our quick tests! Please proceed!")
```

You have passed our quick tests! Please proceed!

2. Complete the function below which should take a `user_id` and provide an ordered list of the most similar users to that user (from most similar to least similar). The returned result should not contain the provided `user_id`, as we know that each user is similar to him/herself. Because the results for each user here are binary, it (perhaps) makes sense to compute similarity as the dot product of two users.

Use the tests to test your function.

```
def find_similar_users(user_id, user_item=user_item):
    """
    INPUT:
    user_id - (int) a user_id
    user_item - (pandas dataframe) matrix of users by articles:
                1's when a user has interacted with an article, 0
    otherwise

    OUTPUT:
    similar_users - (list) an ordered list where the closest users
    (largest dot product users)
                are listed first

    Description:
    Computes the similarity of every pair of users based on the dot
    product
    Returns an ordered

    """
    # compute similarity of each user to the provided user
    similarity_series = user_item.dot(user_item.T)[user_id]
    # sort by similarity
    most_similar_users = similarity_series.sort_values(ascending =
False)
    # create list of just the ids
    most_similar_users = list(most_similar_users.index)
    # remove the own user's id
    most_similar_users.remove(user_id)

    return most_similar_users # return a list of the users in order
    from most to least similar

# Do a spot check of your function
print("The 10 most similar users to user 1 are:
{}".format(find_similar_users(1)[:10]))
print("The 5 most similar users to user 3933 are:
{}".format(find_similar_users(3933)[:5]))
print("The 3 most similar users to user 46 are:
{}".format(find_similar_users(46)[:3]))
```

The 10 most similar users to user 1 are: [3933, 23, 3782, 203, 4459, 3870, 131, 4201, 46, 5041]

The 5 most similar users to user 3933 are: [1, 23, 3782, 203, 4459]

The 3 most similar users to user 46 are: [4201, 3782, 23]

3. Now that you have a function that provides the most similar users to each user, you will want to use these users to find articles you can recommend. Complete the functions below to return the articles you would recommend to each user.

```
df.head()
```

	article_id	title
user_id		
0	1430.0	using pixiedust for fast, flexible, and easier...
1		
1	1314.0	healthcare python streaming application demo
2		
2	1429.0	use deep learning for image classification
3		
3	1338.0	ml optimization using cognitive assistant
4		
4	1276.0	deploy your python model as a restful api
5		

```
user_item.head()
```

article_id	0.0	2.0	4.0	8.0	9.0	12.0	14.0
15.0 \							
user_id							
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0
0.0							
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0
0.0							
3	0.0	0.0	0.0	0.0	0.0	1.0	0.0
0.0							
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0
0.0							
5	0.0	0.0	0.0	0.0	0.0	0.0	0.0
0.0							

article_id	16.0	18.0	...	1434.0	1435.0	1436.0	1437.0
1439.0 \							
user_id			...				
1	0.0	0.0	...	0.0	0.0	1.0	0.0
1.0							
2	0.0	0.0	...	0.0	0.0	0.0	0.0
0.0							
3	0.0	0.0	...	0.0	0.0	1.0	0.0
0.0							

4	0.0	0.0	...	0.0	0.0	0.0	0.0
0.0							
5	0.0	0.0	...	0.0	0.0	0.0	0.0
0.0							

article_id	1440.0	1441.0	1442.0	1443.0	1444.0
user_id					
1	0.0	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0
5	0.0	0.0	0.0	0.0	0.0

[5 rows x 714 columns]

```
def get_article_names(article_ids, df=df):
    """
    INPUT:
    article_ids - (list) a list of article ids
    df - (pandas dataframe) df as defined at the top of the notebook

    OUTPUT:
    article_names - (list) a list of article names associated with the
    list of article ids
    """
    # create a list of article names identified by the the title
    column associated with the list of article ids
    article_names = [df[df["article_id"] == float(x)]["title"].max()
    for x in article_ids]

    return article_names # Return the article names associated with
list of article ids

def get_user_articles(user_id, user_item=user_item):
    """
    INPUT:
    user_id - (int) a user id
    user_item - (pandas dataframe) matrix of users by articles:
    1's when a user has interacted with an article, 0
    otherwise

    OUTPUT:
    article_ids - (list) a list of the article ids seen by the user
    article_names - (list) a list of article names associated with the
    list of article ids
    """
    # (this is identified by the doc_full_name column in
    df_content)
```

```

    Description:
    Provides a list of the article_ids and article titles that have
    been seen by a user
    '''
    # pull all the article ids for the user which donot zeros for
    their values in the user_item df
    article_ids = user_item.loc[user_id][user_item.loc[user_id] !=
0].index.astype(str)
    # pass the article_ids to the get_article_names function to get
    the article names for those ids
    article_names = get_article_names(article_ids)

    return article_ids, article_names # return the ids and names

def user_user_recs(user_id, m=10):
    '''
    INPUT:
    user_id - (int) a user id
    m - (int) the number of recommendations you want for the user

    OUTPUT:
    recs - (list) a list of recommendations for the user

    Description:
    Loops through the users based on closeness to the input user_id
    For each user - finds articles the user hasn't seen before and
    provides them as recs
    Does this until m recommendations are found

    Notes:
    Users who are the same closeness are chosen arbitrarily as the
    'next' user

    For the user where the number of recommended articles starts below
m
    and ends exceeding m, the last items are chosen arbitrarily
    '''
    # Ids of articles seen by user (we don't want to recommend these)
    user_seen = get_user_articles(user_id)[0]
    # Similar users to the current user
    similar_users = find_similar_users(user_id)

    # Keep the recommended articles here
    recs = []

    # Go through the similar users and identify articles they like the
    user hasn't seen
    for similar_user in similar_users:

```



```

        similar_user_seen = get_user_articles(similar_user)[0]
        # Obtain recommendations from similar user without the
        articles seen by the user
        new_recs = np.setdiff1d(similar_user_seen, user_seen,
assume_unique=True)
        # Update recs with new recs; concat rec and new_rec arrays and
        return only unique values
        recs = np.unique(np.concatenate([new_recs, recs], axis = 0))

        # If we have enough recommendations exit the loop
        if len(recs) > m-1:
            break

    return recs[:m] # return your recommendations for this user_id

# Check Results
get_article_names(user_user_recs(1, 10)) # Return 10 recommendations
for user 1

['recommender systems: approaches & algorithms',
'1448    i ranked every intro to data science course on...\nName:
title, dtype: object',
'data tidying in data science experience',
'a tensorflow regression model to predict house values',
'520    using notebooks with pixiedust for fast, flexi...\nName:
title, dtype: object',
'airbnb data for analytics: mallorca reviews',
'airbnb data for analytics: vancouver listings',
'analyze facebook data using ibm watson and watson studio',
'analyze accident reports on amazon emr spark',
'analyze energy consumption in buildings']

# Test your functions here - No need to change this code - just run
this cell
assert set(get_article_names(['1024.0', '1176.0', '1305.0', '1314.0',
'1422.0', '1427.0'])) == set(['using deep learning to reconstruct
high-resolution audio', 'build a python app on the streaming analytics
service', 'gosales transactions for naive bayes model', 'healthcare
python streaming application demo', 'use r dataframes & ibm watson
natural language understanding', 'use xgboost, scikit-learn & ibm
watson machine learning apis']), "Oops! Your the get_article_names
function doesn't work quite how we expect."
assert set(get_article_names(['1320.0', '232.0', '844.0'])) ==
set(['housing (2015): united states demographic measures', 'self-
service data preparation with ibm data refinery', 'use the cloudant-
spark connector in python notebook']), "Oops! Your the
get_article_names function doesn't work quite how we expect."
assert set(get_user_articles(20)[0]) == set(['1320.0', '232.0',
'844.0'])

```

```

assert set(get_user_articles(20)[1]) == set(['housing (2015): united
states demographic measures', 'self-service data preparation with ibm
data refinery','use the cloudbant-spark connector in python notebook'])
assert set(get_user_articles(2)[0]) == set(['1024.0', '1176.0',
'1305.0', '1314.0', '1422.0', '1427.0'])
assert set(get_user_articles(2)[1]) == set(['using deep learning to
reconstruct high-resolution audio', 'build a python app on the
streaming analytics service', 'gosales transactions for naive bayes
model', 'healthcare python streaming application demo', 'use r
dataframes & ibm watson natural language understanding', 'use xgboost,
scikit-learn & ibm watson machine learning apis'])
print("If this is all you see, you passed all of our tests! Nice
job!")

```

If this is all you see, you passed all of our tests! Nice job!

4. Now we are going to improve the consistency of the **user_user_recs** function from above.

- Instead of arbitrarily choosing when we obtain users who are all the same closeness to a given user - choose the users that have the most total article interactions before choosing those with fewer article interactions.
- Instead of arbitrarily choosing articles from the user where the number of recommended articles starts below m and ends exceeding m, choose articles with the articles with the most total interactions before choosing those with fewer total interactions. This ranking should be what would be obtained from the **top_articles** function you wrote earlier.

```

def get_top_sorted_users(user_id, df=df, user_item=user_item):
    """
    INPUT:
    user_id - (int)
    df - (pandas dataframe) df as defined at the top of the notebook
    user_item - (pandas dataframe) matrix of users by articles:
                1's when a user has interacted with an article, 0
    otherwise

    OUTPUT:
    neighbors_df - (pandas dataframe) a dataframe with:
                    neighbor_id - is a neighbor user_id
                    similarity - measure of the similarity of each
user to the provided user_id
                    num_interactions - the number of articles viewed
by the user - if a u

    Other Details - sort the neighbors_df by the similarity and then
by number of interactions where
                    highest of each is higher in the dataframe

```

```

'''
    # create an empty dataframe with "neighbor_id",
    "similarity", "num_interactions" columns
    neighbors_df=pd.DataFrame(columns=["neighbor_id",
    "similarity", "num_interactions"])

    # set the "neighbor_id" column and remove the own user_id of the
    user
    neighbors_df["neighbor_id"]=[id for id in
    range(1,user_item.shape[0]) if id != user_id]

    # set the "similarity" column to most similar using dot product
    neighbors_df["similarity"] =
    neighbors_df["neighbor_id"].apply(lambda x:
    np.dot(user_item.loc[user_id], user_item.loc[x]))

    # set the "num_interactions" column to user_interactions_counts
    series
    # create a series of user_interactions_counts created using
    grouping of "user_id"
    user_interactions_counts = df.groupby(["user_id"])
    ["article_id"].count()

    neighbors_df["num_interactions"]=neighbors_df["neighbor_id"].apply(lambda x: user_interactions_counts.loc[x].sum())

    # sort dataframe ascending by "similarity", "num_interactions" and
    reset index to start from 0

    neighbors_df=neighbors_df.sort_values(by=["similarity", "num_interactions"], ascending=False).reset_index(drop=True)

    return neighbors_df # Return the dataframe specified in the
    doc_string

```

```

def user_user_recs_part2(user_id, m=10):
    '''
    INPUT:
    user_id - (int) a user id
    m - (int) the number of recommendations you want for the user

    OUTPUT:
    recs - (list) a list of recommendations for the user by article id

```

rec_names - (list) a list of recommendations for the user by article title

Description:

*Loops through the users based on closeness to the input user_id
For each user - finds articles the user hasn't seen before and provides them as recs*

Does this until m recommendations are found

Notes:

** Choose the users that have the most total article interactions before choosing those with fewer article interactions.*

** Choose articles with the articles with the most total interactions 4 before choosing those with fewer total interactions.*

'''

article ids,names read by our user ; don't want to recommend these

*user_articles_ids_read, user_articles_names_read =
get_user_articles(user_id, user_item)*

top similar user ids closest to our user

*similar_users = get_top_sorted_users(user_id, df, user_item)
["neighbor_id"]*

keep the recommended articles here

recs = []

for similar_user **in** similar_users:

articles read by similar user

*sim_articles_ids_read, sim_articles_names_read =
get_user_articles(similar_user, user_item)*

obtain recommendations from similar user without the articles read by our user

*new_rec = np.setdiff1d(sim_articles_ids_read,
user_articles_ids_read, assume_unique=True)*

Update recs with new recs; concat rec and new_rec arrays and return only unique values

recs = np.unique(np.concatenate([new_rec, recs], axis = 0))

break the loop once the number of recs have been reached

if len(recs) > m-1:

break

recs = recs[:m]

#recs = recs.tolist() # convert to a list

```
rec_names = get_article_names(recs)
```

```
return recs, rec_names
```

```
# Quick spot check - don't change this code - just use it to test your functions
```

```
rec_ids, rec_names = user_user_recs_part2(20, 10)
print("The top 10 recommendations for user 20 are the following article ids:")
print(rec_ids)
print()
print("The top 10 recommendations for user 20 are the following article names:")
print(rec_names)
```

```
The top 10 recommendations for user 20 are the following article ids:
['1024.0' '1085.0' '109.0' '1150.0' '1151.0' '1152.0' '1153.0'
'1154.0'
'1157.0' '1160.0']
```

```
The top 10 recommendations for user 20 are the following article names:
```

```
['using deep learning to reconstruct high-resolution audio', 'airbnb data for analytics: chicago listings', 'tensorflow quick tips', 'airbnb data for analytics: venice calendar', 'airbnb data for analytics: venice listings', 'airbnb data for analytics: venice reviews', 'airbnb data for analytics: vienna calendar', 'airbnb data for analytics: vienna listings', 'airbnb data for analytics: washington d.c. listings', 'analyze accident reports on amazon emr spark']
```

5. Use your functions from above to correctly fill in the solutions to the dictionary below. Then test your dictionary against the solution. Provide the code you need to answer each following the comments below.

```
get_top_sorted_users(1)
```

	neighbor_id	similarity	num_interactions
0	3933	35.0	45
1	23	17.0	364
2	3782	17.0	363
3	203	15.0	160
4	4459	15.0	158
...
5142	5131	0.0	1
5143	5141	0.0	1
5144	5144	0.0	1
5145	5147	0.0	1
5146	5148	0.0	1

```
[5147 rows x 3 columns]
```

```
get_top_sorted_users(1).iloc[0]
```

```
neighbor_id      3933.0
similarity        35.0
num_interactions  45.0
Name: 0, dtype: float64
```

```
get_top_sorted_users(131).iloc[9]
```

```
neighbor_id      242.0
similarity        25.0
num_interactions 148.0
Name: 9, dtype: float64
```

```
### Tests with a dictionary of results
```

```
user1_most_sim = get_top_sorted_users(1).iloc[0].neighbor_id # Find
the user that is most similar to user 1
user131_10th_sim = get_top_sorted_users(131).iloc[9].neighbor_id #
Find the 10th most similar user to user 131
```

```
## Dictionary Test Here
```

```
sol_5_dict = {
    'The user that is most similar to user 1.': user1_most_sim,
    'The user that is the 10th most similar to user 131':
    user131_10th_sim,
}
```

```
t.sol_5_test(sol_5_dict)
```

This all looks good! Nice job!

6. If we were given a new user, which of the above functions would you be able to use to make recommendations? Explain. Can you think of a better way we might make recommendations? Use the cell below to explain a better method for new users.

Provide your response here.

If we were given a new user, I would use the `get_top_articles()` function and provide the user with the rank based recommendations (suggest the top ranked articles) and that is because we do not have any user-article interaction information about the user. We cannot find similar users (neighbours) if we do not have that information. The other methods individually (content, collaborative, knowledge) may not be as useful in such a situation.

7. Using your existing functions, provide the top 10 recommended articles you would provide for the a new user below. You can test your function against our thoughts to make sure we are all on the same page with how we might make a recommendation.

```
new_user = '0.0'
```

```
# What would your recommendations be for this new user '0.0'? As a
new user, they have no observed articles.
```

```
# Provide a list of the top 10 article ids you would give to
new_user_recs = get_top_article_ids(10)# Your recommendations here
```

new_user_recs

```
['1429.0',
 '1330.0',
 '1431.0',
 '1427.0',
 '1364.0',
 '1314.0',
 '1293.0',
 '1170.0',
 '1162.0',
 '1304.0']
```

```
assert set(new_user_recs) ==
set(['1314.0', '1429.0', '1293.0', '1427.0', '1162.0', '1364.0', '1304.0', '1
170.0', '1431.0', '1330.0']), "Oops! It makes sense that in this case
we would want to recommend the most popular articles, because we don't
know anything about these users."
```

```
print("That's right! Nice job!")
```

That's right! Nice job!

Part V: Matrix Factorization

In this part of the notebook, you will build use matrix factorization to make article recommendations to the users on the IBM Watson Studio platform.

1. You should have already created a **user_item** matrix above in **question 1** of **Part III** above. This first question here will just require that you run the cells to get things set up for the rest of **Part V** of the notebook.

```
# Load the matrix here
```

```
user_item_matrix = pd.read_pickle('user_item_matrix.p')
```

```
# a quick look at the matrix
```

```
user item matrix.head()
```

```

article_id  0.0   100.0  1000.0  1004.0  1006.0  1008.0  101.0  1014.0
1015.0 \
user id

```

1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
0.0								
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

```

0.0
3      0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0
0.0
4      0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0
0.0
5      0.0    0.0    0.0    0.0    0.0    0.0    0.0    0.0
0.0

```

```

article_id 1016.0 ... 977.0 98.0 981.0 984.0 985.0 986.0
990.0 \
user_id      ...

```

```

1      0.0 ... 0.0 0.0 1.0 0.0 0.0 0.0
0.0
2      0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0
0.0
3      0.0 ... 1.0 0.0 0.0 0.0 0.0 0.0
0.0
4      0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0
0.0
5      0.0 ... 0.0 0.0 0.0 0.0 0.0 0.0
0.0

```

```

article_id 993.0 996.0 997.0
user_id
1      0.0    0.0    0.0
2      0.0    0.0    0.0
3      0.0    0.0    0.0
4      0.0    0.0    0.0
5      0.0    0.0    0.0

```

[5 rows x 714 columns]

2. In this situation, you can use Singular Value Decomposition from [numpy](#) on the user-item matrix. Use the cell to perform SVD, and explain why this is different than in the lesson.

Perform SVD on the User-Item Matrix Here

u, s, vt = np.linalg.svd(user_item_matrix)# use the built in to get the three matrices

checking to see if any rows in the user_item_matrix has an NaNs
user_item_matrix.isna().sum().sort_values(ascending=False)

```

article_id
0.0      0
485.0    0
443.0    0
444.0    0
446.0    0

```



```

1338.0    ..
134.0     0
1340.0    0
1343.0    0
997.0     0
Length: 714, dtype: int64

```

there are no rows in the matrix that has null values

checking to see if any rows in the user_item_matrix have values other than 0 and 1

```
(~user_item_matrix.isin([0,1])).sum().sort_values(ascending=False)
```

```

article_id
0.0      0
485.0    0
443.0    0
444.0    0
446.0    0
..
1338.0    0
134.0     0
1340.0    0
1343.0    0
997.0     0
Length: 714, dtype: int64

```

fortunately there are no rows in the matrix that have values other than 0's and 1's

Provide your response here.

As there are no null values in the user_item_matrix, SVD is a very good candidate. Also the matrix has only 1's and 0's (interactions vs no interactions). As far as the lesson is concerned the matrix had lot of null values and FunkSVD was opted. But in our scenario SVD should work just fine.

3. Now for the tricky part, how do we choose the number of latent features to use? Running the below cell, you can see that as the number of latent features increases, we obtain a lower error rate on making predictions for the 1 and 0 values in the user-item matrix. Run the cell below to get an idea of how the accuracy improves as we increase the number of latent features.

```

num_latent_feats = np.arange(10,700+10,20)
sum_errs = []

```

```

for k in num_latent_feats:
    # restructure with k latent features
    s_new, u_new, vt_new = np.diag(s[:k]), u[:, :k], vt[:k, :]

    # take dot product

```

```


user_item_est = np.around(np.dot(np.dot(u_new, s_new), vt_new))

# compute error for each prediction to actual value
diffs = np.subtract(user_item_matrix, user_item_est)

# total errors and keep track of them
err = np.sum(np.sum(np.abs(diffs)))
sum_errs.append(err)

plt.plot(num_latent_feats, 1 - np.array(sum_errs)/df.shape[0]);
plt.xlabel('Number of Latent Features');
plt.ylabel('Accuracy');
plt.title('Accuracy vs. Number of Latent Features');

```



4. From the above, we can't really be sure how many features to use, because simply having a better way to predict the 1's and 0's of the matrix doesn't exactly give us an indication of if we are able to make good recommendations. Instead, we might split our dataset into a training and test set of data, as shown in the cell below.

Use the code from question 3 to understand the impact on accuracy of the training and test sets of data with different numbers of latent features. Using the split below:

- How many users can we make predictions for in the test set?
- How many users are we not able to make predictions for because of the cold start problem?

- How many articles can we make predictions for in the test set?
- How many articles are we not able to make predictions for because of the cold start problem?

```
df_train = df.head(40000)
```

```
df_test = df.tail(5993)
```

```
def create_test_and_train_user_item(df_train, df_test):
```

```
    '''
```

```
    INPUT:
```

```
    df_train - training dataframe
```

```
    df_test - test dataframe
```

```
    OUTPUT:
```

```
    user_item_train - a user-item matrix of the training dataframe  
                     (unique users for each row and unique articles
```

```
for each column)
```

```
    user_item_test - a user-item matrix of the testing dataframe  
                   (unique users for each row and unique articles for
```

```
each column)
```

```
    test_idx - all of the test user ids
```

```
    test_arts - all of the test article ids
```

```
    '''
```

```
    user_item_train = create_user_item_matrix(df_train) # Returns a  
train matrix with user ids as rows and article ids on the columns
```

```
    user_item_test = create_user_item_matrix(df_test) # Returns a test  
matrix with user ids as rows and article ids on the columns
```

```
    test_idx = user_item_test.index # all of the test user ids
```

```
    test_arts = user_item_test.columns # all of the test article ids
```

```
    return user_item_train, user_item_test, test_idx, test_arts
```

```
user_item_train, user_item_test, test_idx, test_arts =
```

```
create_test_and_train_user_item(df_train, df_test)
```

```
# How many users can we make predictions for in the test set?
```

```
test_set_users = len(test_idx)
```

```
print("No of users in test set: ", test_set_users) # no. of users in  
test set
```

```
train_idx = user_item_train.index # users in training set
```

```
print("No of users in train set: ", len(train_idx))
```

```
train_test_user_intersect = len((test_idx).intersection(train_idx)) #  
Return the number of users exist in both test and train
```

```
print("No of users we can make predictions for in train set:
```

```
", train_test_user_intersect)
```

```
# How many users in the test set are we not able to make predictions
for because of the cold start problem?
print("No of users in the test set are we not able to make predictions
for because of the cold start problem: ",
np.abs((train_test_user_intersect)-len(test_idx)))
```

```
# How many movies can we make predictions for in the test set?
print("No of articles can we make predictions for in the test set:
",len(test_arts))
```

```
# How many articles are we not able to make predictions for because of
the cold start problem?
print("No of articles in the test set: ",len(test_arts))
train_arts=user_item_train.columns
print("No of articles in the train set: ",len(train_arts))
```

```
train_test_articles_diff = len((test_arts).difference(train_arts)) #
Return the number of articles exist in test but not train
print("Movies in the test set we're not able to make predictions for
because of the cold start problem: ",train_test_articles_diff)
```

```
No of users in test set: 682
No of users in train set: 4487
No of users we can make predictions for in train set: 20
No of users in the test set are we not able to make predictions for
because of the cold start problem: 662
No of articles can we make predictions for in the test set: 574
No of articles in the test set: 574
No of articles in the train set: 714
Movies in the test set we're not able to make predictions for because
of the cold start problem: 0
```

```
# Replace the values in the dictionary below
a = 662
b = 574
c = 20
d = 0
```

```
sol_4_dict = {
    'How many users can we make predictions for in the test set?': c,
    'How many users in the test set are we not able to make
predictions for because of the cold start problem?': a,
    'How many movies can we make predictions for in the test set?': b,
    # replacing "articles" with "movies" to match code in project_tests.py
https://knowledge.udacity.com/questions/770708
    'How many movies in the test set are we not able to make
predictions for because of the cold start problem?': d # replacing
"articles" with "movies" to match code in project_tests.py
https://knowledge.udacity.com/questions/770708
```

```
}
```

```
t.sol_4_test(sol_4_dict)
```

Awesome job! That's right! All of the test movies are in the training data, but there are only 20 test users that were also in the training set. All of the other users that are in the test set we have no data on. Therefore, we cannot make predictions for these users using SVD.

5. Now use the **user_item_train** dataset from above to find U, S, and V transpose using SVD. Then find the subset of rows in the **user_item_test** dataset that you can predict using this matrix decomposition with different numbers of latent features to see how many features makes sense to keep based on the accuracy on the test data. This will require combining what was done in questions 2 - 4.

Use the cells below to explore how well SVD works towards making predictions for recommendations on the test data.

```
# lets look at a sample of 3 records from user_item_train  
user_item_train.sample(3)
```

```
article_id  0.0      2.0      4.0      8.0      9.0      12.0      14.0  
15.0 \  
user_id  
  
590          0.0      0.0      0.0      0.0      0.0      0.0      0.0  
0.0  
2756         0.0      0.0      0.0      0.0      0.0      0.0      0.0  
0.0  
3219         0.0      0.0      0.0      0.0      0.0      0.0      0.0  
0.0
```

```
article_id  16.0     18.0     ...  1434.0  1435.0  1436.0  1437.0  
1439.0 \  
user_id          ...  
  
590          0.0      0.0     ...      0.0      0.0      0.0      0.0  
0.0  
2756         0.0      0.0     ...      0.0      1.0      0.0      0.0  
0.0  
3219         0.0      0.0     ...      0.0      0.0      0.0      0.0  
0.0
```

```
article_id  1440.0  1441.0  1442.0  1443.0  1444.0  
user_id  
590          0.0      0.0      0.0      0.0      0.0  
2756         0.0      0.0      0.0      0.0      0.0  
3219         0.0      0.0      0.0      0.0      0.0
```

[3 rows x 714 columns]

fit SVD on the user_item_train matrix

*u_train, s_train, vt_train = np.linalg.svd(user_item_train)# fit svd
similar to above then use the cells below*

print the shapes of u,s,vt_train

*print(f"u_train_shape: {u_train.shape}, s_train_shape:
{s_train.shape}, vt_train_shape: {vt_train.shape} ")*

*u_train_shape: (4487, 4487), s_train_shape: (714,), vt_train_shape:
(714, 714)*

*# Use these cells to see how well you can use the training
decomposition to predict on test data*

latent features

num_latent_feats = np.arange(10,700+10,20)

keep sum of errors in train here

sum_errors_train = []

keep sum of errors in test here

sum_errors_test = []

common users

row_index = user_item_train.index.isin(test_idx)

common articles

col_index = user_item_train.columns.isin(test_arts)

u_test = u_train[row_index, :]

vt_test = vt_train[:, col_index]

users we can predict for in test

*test_users_predict =
np.intersect1d(user_item_train.index.tolist(),user_item_test.index.to
list())*

test matrix with the users from above

user_intersect_train_test = user_item_test.loc[test_users_predict]

for k **in** num_latent_feats:

restructure with k latent features

*s_train_new, u_train_new, vt_train_new = np.diag(s_train[:k]),
u_train[:, :k], vt_train[:k, :]
u_test_new, vt_test_new = u_test[:, :k], vt_test[:k, :]*

take dot product

*user_item_train_pred = np.around(np.dot(np.dot(u_train_new,
s_train_new), vt_train_new))*

*user_item_test_pred = np.around(np.dot(np.dot(u_test_new,
s_train_new), vt_test_new))*

```

    # compute error for each prediction to actual value
    diffs_train = np.subtract(user_item_train, user_item_train_pred)
    diffs_test = np.subtract(user_intersect_train_test,
user_item_test_pred)

    # calculate total errors
    train_errors = np.sum(np.sum(np.abs(diffs_train)))
    test_errors = np.sum(np.sum(np.abs(diffs_test)))

    # keep track of total errors
    sum_errors_train.append(train_errors)
    sum_errors_test.append(test_errors)


# plot the train and test accuracy vs. the number of latent features
fig, ax = plt.subplots(figsize = (10, 6))

plt.plot(num_latent_feats,
1 -
(np.array(sum_errors_train)/(user_item_train.shape[0]*user_item_train.
shape[1])),
label="Train");

plt.plot(num_latent_feats,
1 -
(np.array(sum_errors_test)/(user_intersect_train_test.shape[0]*user_in
tersect_train_test.shape[1])),
label="Test");

ax.grid("on")
plt.legend();
plt.xlabel("Number of Latent Features");
plt.ylabel("Accuracy");
plt.title("Accuracy vs. Number of Latent Features");

```



6. Use the cell below to comment on the results you found in the previous question. Given the circumstances of your results, discuss what you might do to determine if the recommendations you make with any of the above recommendation systems are an improvement to how users currently find articles?

Your response here.

As we can see from the visualization the accuracy of the training set increases as the number of latent features increases, but sadly that's not the case with our test set. For the test data with increase of latent features the accuracy only tends to decrease. One plausible cause is overfitting meaning the model performs very well on the training data but does not generalize well on test and inference data (<https://stackoverflow.com/questions/37776333/why-too-many-features-cause-overfitting>). To avoid this scenario we may need to lower the number of latent features.

Also we have noticed that the matrix used is loaded with sparse features which might also be causing the model not to perform well on the training data (<https://www.kdnuggets.com/2021/01/sparse-features-machine-learning-models.html#:~:text=Model%20algorithms%20and%20diagnostic%20measures,This%20is%20called%20overfitting>).

We know that we can only make predictions for 20 users in the training data as there is only a small overlap of users between these splits. Maybe an improved accuracy can be observed with an increase in the overlap of users and not just 20 users or we can use a rank based recommendation system for the rest of users that we cannot make recommendations for. Also another method to consider would be to create an A/B test or something similar to it to measure how we can use our new recommendation system to a section of users and rank based recommendation system to the other group and measure

the click rate and time spent on the recommended articles and determine if our new recommendation system is working. When we look at training data accuracy in the visual above there is a steep drop after 50 features and continues to drop till it plateaued out at 400 features from where it remained flat. So around 50 latent features would be an optimal number of latent features in this case.

References

<https://www.shanelynn.ie/bar-plots-in-python-using-pandas-dataframes/>
<https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.duplicated.html>
<https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.unstack.html>
<https://towardsdatascience.com/introduction-to-pandas-apply-applymap-and-map-5d3e044e93ff>
<https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.transpose.html>
<https://stackoverflow.com/questions/27967914/efficient-way-to-compute-intersecting-values-between-two-numpy-> <https://knowledge.udacity.com/questions/770708>
<https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.Index.difference.html> <https://pandas.pydata.org/docs/reference/api/pandas.Index.intersection.html>
https://matplotlib.org/stable/api/_as_gen/matplotlib.pyplot.legend.html
<https://stackoverflow.com/questions/27890735/difference-between-listnumpy-array-and-numpy-array-tolist>